

A PRIVATE DIGEST · NO. 2

The Elliot Edition

The week the AI investment thesis met actual economics: hidden trillions, a fund collapse, and the quiet practitioners who kept building anyway.

JULY 26 – AUGUST 2, 2026

EIGHT STORIES

In This Issue

– Editor’s Note

01 Big Tech’s Hidden AI Debt & the Chips Meltdown

02 The Iran War Drags On: Politics and the Bond Market

03 The Open vs Closed AI Debate

04 Autonomous Vehicles: The \$11 Trillion Transition

05 Google’s Record Quarter & the Space Economy

06 Markets & Investing

07 AI Tools, Claude, and the Practitioner Edge

08 Creator Economy & Media

– *Week in One Thought*

The Week the Numbers Arrived

The week of July 26 through August 2 was one of those moments when several distinct storylines converged into a single, uncomfortable diagnosis. The hidden debt of the hyperscalers, the chips meltdown, the geopolitical premium in bond markets, the ideological battle over open versus closed AI—these threads are separate only in how they are covered. Underneath, they describe the same thing: a technology cycle that has moved faster than the financial and regulatory infrastructure built to contain it.

What gave the week its peculiar texture was the contrast between the headlines and what was happening in practitioners' inboxes. While markets were processing the news that a major AI hedge fund had lost sixty-seven percent in a single month, the people actually building with AI were publishing some of the most operationally useful guidance the field has seen—meticulous analyses of token costs, context window management, and model selection. The gap between the financial abstraction and the engineering reality has rarely felt wider.

The gap between the financial abstraction and the engineering reality has rarely felt wider—and that gap is where this issue lives.

The geopolitical subtext running through everything was the Iran conflict, now approaching six months. It has done something subtle to the bond market's psychology: the thirty-year Treasury yield at 5.23 percent, its

highest since 2007, is not just a number about interest rates. It is a number about how much borrowing room remains for an AI buildout that Goldman Sachs models at \$5.3 trillion

in combined hyperscaler capex through 2030.

What follows across eight sections is an attempt to read all of it together. The instinct to separate the geopolitics from the markets from the technology

is understandable—but it produces a false picture. The people who navigated the last two technology cycles with their capital and judgment intact were the ones who refused to compartmentalise.

OFF THE BOOKS

Hidden Trillions & the Chips Meltdown

How \$1.65 trillion in off-balance-sheet debt, a Korean market rout, and a hedge fund blowup put the entire AI investment thesis under stress in a single week.

\$1.65T

When the Bubble Meets the Balance Sheet

A Nikkei investigation, a Korean market rout, and a hedge fund fire sale converged to price in a new question: who actually bears the risk?

The week's defining financial story opened with a number buried in a **Nikkei Asia** investigation: \$1.65 trillion in off-balance-sheet debt quietly accumulated by the five largest hyperscalers—Alphabet, Amazon, Meta, Microsoft, and Oracle. The mechanism is relatively straightforward but its scale is staggering. Instead of building data centres on their own balance sheets, these companies sign long-term lease agreements with operators—operators that are often shell companies funded by private credit providers including **Blue Owl Capital**,

Apollo Global, and **Blackstone**. Those funds are in turn financed by pension funds and life insurance companies.

Scott Galloway and co-writer Ed Elson noted that the California State Teachers' Retirement System is among Blue Owl's largest investors. The implication is sharp: if the AI bubble bursts, the risk does not fall on Big Tech's balance sheets. It falls on the retirement savings of working people, while remaining invisible to the SEC and Federal Reserve, who have no regulatory jurisdiction over private credit.

“

The risk does not fall on Big Tech's balance sheets. It falls on the retirement savings of ordinary working people.

— SCOTT GALLOWAY, PROF G MARKETS

On Tuesday, **DealBook's** "Chips Meltdown" edition arrived with its own alarming figures. South Korea's Kospi index fell more than ten percent in a single session, wiping nearly \$600 billion from **SK Hynix's** market cap. **Nvidia** tumbled five percent in the US, briefly surrendering its crown as the world's most valuable listed company to Apple. **Andrew Ross Sorkin** identified three overlapping forces: anxiety about whether AI capex is generating real returns; concerns about circular financing (Nvidia is reportedly discussing a \$250 billion guarantee for **OpenAI's** Ohio data centre); and growing competitive pressure from China. **Goldman Sachs** projects the four biggest hyperscalers will spend a combined \$5.3 trillion on capital expenditure through 2030 and plan to borrow up to \$400 billion this year alone.

By Wednesday, **Microsoft** reported a 31% quarterly profit jump to \$35.8 billion, with cloud revenue growing at its fastest rate in four years. **Meta** fell sharply in premarket. The contrast crystallised investor anxiety: Microsoft has demonstrable AI revenue; Meta is still spending without the same evidence of return.

The sharpest analytical framing came from **GTMnow's** Friday piece on the token price collapse. Token prices have fallen more than 95% in three years—from \$30 per million input tokens in

March 2023 to under \$0.50 today. Yet enterprise LLM spend more than doubled in six months, from \$3.5 billion to \$8.4 billion. **Sundar Pichai** revealed Google went from processing 9.7 trillion to 3.2 quadrillion monthly tokens—a 330-fold jump—in the same period that per-token prices fell continuously. GTMnow's framing draws on the Jevons Paradox: cheap tokens drive more token use, not less. The structural consequence for AI-native businesses: 50–60% gross margins rather than the 80–90% typical of traditional SaaS, because every inference call is a variable cost of goods sold. One cautionary tale surfaced: a four-agent system with no hard cost cap looped for eleven days and burned \$47,000 before anyone caught it.

Then Friday brought the week's most dramatic coda. **DealBook** reported that Situational Awareness, the AI hedge fund founded two years ago by former **OpenAI** researcher Leopold Aschenbrenner, had seen its portfolio fall roughly 67% in a single month, forcing a fire sale to **Ken Griffin's** Citadel to cover margin calls. The fund had raised over \$10 billion—backed by **Stripe's** Patrick and John Collison and former **GitHub** CEO Nat Friedman—held a \$3.5 billion Anthropic stake, and been up roughly 80% on the year before the collapse. Roger Lowenstein, author of *When Genius Failed*, told DealBook it felt "a little more like the dot-coms—

huge equity investments in a new thing
that no one can value with any hope of
precision.”

SIX MONTHS IN

The Iran War Drags On

Six months into the US-Iran conflict, Trump's approval hits a record low and the thirty-year Treasury yield reaches its highest level since 2007.

66%

SOURCES: LAURA ROZEN / DIPLOMATIC · DEALBOOK · PROF G MARKETS

TRUMP DISAPPROVAL RATING (CNN)

Credibility Shock: When Geopolitics Hits the Bond Market

A record disapproval rating, a dissent-filled Fed vote, and a thirty-year yield not seen since 2007 describe the same underlying fracture.

Six months into the US-Iran conflict, the political and financial costs are becoming impossible to separate. **Laura Rozen**'s Diplomatic newsletter delivered the clearest political read: President Trump hit an all-time high disapproval rating of 66% in a new **CNN** poll, with **Quinnipiac** showing 32% approval and 58% disapproval—Trump's lowest Quinnipiac number on record. More significant is the erosion within his own coalition: Republican approval has slid from 85% in June to 76% now.

Democratic consultant **David Doak** told Rozen that what is unusual about this

moment is not merely that voters are unhappy, but that they have done three specific things: assigned blame to Trump for causing the problems, diagnosed a lack of strategy in Iran, and concluded—most damaging of all—that “he doesn’t care.” Doak argues this trifecta—blame attribution, incompetence diagnosis, perceived indifference—is the hardest combination to recover from politically. Each element alone is survivable; all three together rarely are.

“

Every 0.1 percentage point rise in rates over the coming decade adds \$1.8 trillion to US interest costs—against a national debt already at \$39.4 trillion.

— SCOTT GALLOWAY, PROF G MARKETS

DealBook's Thursday "Credibility Shock" piece connected the war directly to the fixed-income market. The 30-year Treasury yield rose to 5.23%—its highest since 2007, right before the global financial crisis. Traders priced in 60% odds of a September rate increase.

Fed Chairman Kevin Warsh held rates steady in a dissent-heavy 9-to-3 vote but sent mixed signals: he opened the door to measuring inflation against indicators beyond PCE, suggested markets had already done some tightening for the Fed, and simultaneously insisted the Fed would not go "soft" on its 2% target. **Bank of America** economist Aditya Bhawe characterised the market moves as

"consistent with a central bank inflation credibility shock."

The fiscal vulnerability beneath all of this has been quantified by Galloway: every 0.1 percentage point rise in rates over the coming decade adds approximately \$1.8 trillion to US interest costs, against a national debt already at \$39.4 trillion. Interest payments on that debt already cost nearly \$1 trillion in fiscal year 2025—approximately \$7,700 per American household—and are on track to become the second-largest federal expenditure in 2026. The arithmetic of wartime deficit spending meeting a yield market repricing is not comfortable.

IDEOLOGY MEETS INFRASTRUCTURE

The Open vs Closed AI Debate

Jensen Huang's 60-million-view post, Zuckerberg's blunt critique, and Chamath's diagnosis of "valuation preservation" opened a new front in AI's ideological war.

60M

Valuation Preservation: The Real Stakes Behind the Open AI Battle

A 60-million-view post, a distillation controversy, and a KYC counterargument reveal that the open-versus-closed debate is as much about moat protection as it is about safety.

The ideological battle that has been simmering in AI for three years burst into public view on Monday when **Jensen Huang** posted on X to announce the Open Secure AI Alliance—a coalition including **Andreessen Horowitz**, **Meta**, and **Microsoft**—arguing that open models would accelerate innovation and strengthen security. The post attracted more than 60 million views. Huang’s framing was diplomatically inclusive: “The world needs both frontier closed models and frontier open models.” The subtext was less so.

As **DealBook** noted, Nvidia has financial positions in both **Anthropic** and **OpenAI**, making its CEO’s public embrace of open models a strategically deft move. Analysts suggested Huang’s logic is that easier access to open-source AI speeds the technology’s diffusion and drives chip demand regardless of which labs prevail. **Mark Zuckerberg** was blunter, telling the *Times* that the argument for keeping powerful models in closed environments was “bad for innovation and even dangerous,” and that AI discourse was “overwhelmingly filled with doom.”

“

What you’re watching is valuation preservation.

— CHAMATH PALIHAPITIYA

Chamath Palihapitiya took the same side but with a more cynical reading. The labs claim that Chinese competitors like **DeepSeek** and **Moonshot AI** are stealing their technology through distillation—running production models at scale to harvest question-and-answer pairs for training. Chamath’s counter: if this were a genuine security concern rather than competitive theatre, the labs would solve it with KYC verification—real identities, bounded credit cards—which would kill industrial-scale account farms immediately. They choose not to, he argues, because it would slow

revenue. His conclusion: “What you’re watching is valuation preservation.”

His broader map of where real value lies: not in the model layer, which is commoditising faster than anyone expected, but above it (applications that people pay for) and below it (chips, cloud, infrastructure). **Treasury Secretary Scott Bessent** threatened sanctions against Chinese companies that “cross the line.” Anthropic and OpenAI formally accused Chinese startups of IP theft via distillation—a charge that, Chamath argues, is as much about market positioning as it is about intellectual property.

THE PRICE OF A MILE

Autonomous Vehicles: \$11 Trillion

From \$1.10 to \$0.25 per mile: Chamath's 92-page Social Capital report sets out the economic thesis that will restructure how a nation moves.

\$11T

From \$1.10 to \$0.25: The Arithmetic of a Transport Revolution

Waymo's 500,000 weekly rides, Aurora's driverless trucks, and an ARK Invest market model point toward a transition that will touch car insurance, personal injury law, and municipal finance.

Chamath Palihapitiya's Friday deep dive on autonomous driving—a 92-page **Social Capital** research report—was the week's most ambitious single piece of analysis. The stakes, as he sets them up: the average American worker loses approximately 225 hours per year to commuting. In 2025, there were over 36,000 traffic fatalities on US roads. **Waymo's** peer-reviewed safety data across 170 million driverless miles shows 92% fewer serious crashes or

fatal injuries compared with human drivers in equivalent conditions.

Waymo passed 20 million cumulative paid driverless rides by end of 2025 and now runs 500,000 rides per week across 28 cities in five countries. The commercial availability of fully autonomous ride-hailing at this scale is, Chamath argues, no longer a question of technology readiness—it is a question of economic velocity.

“

The ‘price of a mile’ is the single metric that unlocks eleven trillion dollars.

— CHAMATH PALIHAPITIYA, SOCIAL CAPITAL

The economic thesis hinges on one metric Chamath calls “the price of a mile.” A personally owned vehicle costs

approximately \$1.10 per mile once you account for insurance, maintenance, depreciation, and fuel. **ARK Invest**

estimates the robotaxi market at \$11 trillion once cost per mile hits \$0.25.

Goldman Sachs expects 65% of US car sales could be autonomous vehicles by 2040. **Aurora** plans 200 or more driverless trucks on the road by end of 2026, targeting the \$900 billion US trucking industry where labour accounts for 45% of per-mile costs.

The second-order disruptions are where Chamath's analysis gets most interesting. Car insurance—a \$330 billion US industry built on the premise of human error—faces structural rethink when 92% of that error is removed from the road. Personal injury law, whose multi-billion-dollar plaintiff bar grew up around car accidents, faces the same question. Speed cameras,

red-light tickets, and parking enforcement collectively generate over \$10 billion annually for municipal governments. Autonomous vehicles don't speed, don't run lights, and navigate to legal parking automatically.

The competitive landscape is settling around three models: fully autonomous ride-hailing (Waymo), semi-autonomous personal vehicles (**Tesla**), and autonomous long-haul trucking (Aurora, **Kodiak**). Chamath's bet is that ride-hailing and trucking will scale fastest because the economics are simplest—a fleet operator can justify the technology on pure cost-per-mile arithmetic without navigating the regulatory complexity of direct consumer sales.

330x AND COUNTING

Google's Record Quarter & the Space Grab

A 330-fold token growth proves the Jevons Paradox in real time, while SpaceX races to claim the invisible real estate that will define the next internet.

330x

GOOGLE TOKEN GROWTH IN 2 YEARS

SOURCES: CHAMATH PALIHAPITIYA · DEALBOOK / PETER COY · ARK INVEST

The Jevons Machine and the Spectrum Land Rush

Google's 330-fold token growth and SpaceX's \$1.8 trillion spectrum bet describe two versions of the same dynamic: when access gets cheaper, demand explodes.

Chamath Palihapitiya's Sunday piece used **Google's** record quarter as the launch pad for a broader argument about AI commoditisation. The headline number from **Sundar Pichai's Google I/O 2026** keynote: token processing grew from 9.7 trillion per month to 3.2 quadrillion—a 330-fold increase—in the same period that per-token prices fell continuously. This is the Jevons Paradox at its clearest: when a resource becomes cheaper, total consumption rises rather than falls. Google is the most visible proof that

cheap tokens do not reduce demand for tokens; they create it.

The implication Chamath draws connects back to the week's broader AI economics story: if token consumption is growing 330-fold while prices collapse, the business model question is not about the model layer at all. It is about the infrastructure layer below it and the application layer above it—the same map he used to argue that the open-versus-closed debate is misdirected.

“

The spectrum wars of the 2020s are the land rushes of the nineteenth century—invisible real estate with trillion-dollar consequences.

— PETER COY, DEALBOOK

The space economy received its own dedicated treatment in **DealBook's** Saturday "Space Grab" piece, written by **Peter Coy**. **SpaceX's** IPO prospectus—which valued the company at \$1.8 trillion—flags spectrum acquisition as a major strategic risk. The invisible real estate that carries wireless signals is being contested in a quiet land rush. **China** has been flooding the **International Telecommunication Union** with spectrum applications while simultaneously advocating for "equitable access"—a defensive positioning against **Starlink's** first-mover advantage.

The coming battle is the World Radiocommunication Conference in

Shanghai in October–November 2027. Coy compares it to the Berlin Conference of 1885: a formal gathering where the real decisions about who gets what will be made in the margins and the bilateral side conversations. Meanwhile, **ARK Invest's** Venture Fund announced an investment in **K2 Space's** Series D—a vertically integrated satellite manufacturer building high-power spacecraft for commercial and national security missions. ARK's Q2 2026 update framed the public listings of SpaceX and **X-energy** as an "inflection point" for the IPO market reopening to category-defining innovation companies.

WAITER OR OWNER

Markets & Investing

GMO challenges the 60/40 portfolio, MastersInvest makes the case for genuine ownership alignment, and the bond market prices in a wartime premium.

5.23%

30-YR TREASURY YIELD (HIGHEST SINCE 2007)

SOURCES: GMO • MASTERSINVEST • DEALBOOK

Three Angles on Portfolio Strategy Under Unusual Stress

GMO's liquid alternatives revival, an employee ownership case study, and a bond market pricing in wartime premium all point in the same direction: alignment matters more than allocation.

GMO published its latest seven-year asset class forecast on Monday alongside a Wednesday piece arguing the case for what they call “the liquid alternatives revival.” The argument challenges how most investors think about diversification. The traditional 60/40 portfolio relies on bonds to hedge equity risk and illiquid alternatives for return potential. GMO’s asset allocation team contends that bonds no longer reliably diversify in every scenario—particularly in an inflationary, wartime environment like the current one, where geopolitical premium is being priced simultaneously

into yields and equities. They further argue that the macro conditions that supported private equity outperformance over the past two decades have ended. Illiquid alternatives, they contend, have been “failing to live up to their inflated promises.”

Their conclusion: liquid alternatives—hedge fund-like strategies accessible in a liquid fund format—may offer the combination of diversification, liquidity, and return potential that the 60/40 model cannot provide in the current regime.

“

Genuine ownership alignment produces businesses that sustainably outperform, because the people closest to the customer have a financial stake in what happens.

The investment thesis is explicit: genuine ownership alignment—where employees receive meaningful economic stakes and the authority to act on them, not motivational language about ownership culture—produces businesses that sustainably outperform because the people with the most

customer contact have the most at stake. Against the backdrop of a week where a \$10 billion AI fund collapsed on leverage and a bond market priced in wartime premium, the case for businesses where alignment is structural rather than notional felt unusually timely.

JULY 26 - AUGUST 2, 2026

1,800 HOURS LATER

AI Tools, Claude, and the Practitioner Edge

Ruben Hassid, Nate's three-part arc, and Lenny Rachitsky's Dianne Penn interview describe a widening gap between those extracting real leverage and those still treating AI like a search engine.

95.73%

REUSED INPUT TOKENS ACROSS 143 THREADS

SOURCES: RUBEN HASSID · NATE · LENNY RACHITSKY

The Widening Gap Between Heavy Users and Everyone Else

Three practitioners—Hassid, Nate, and Rachitsky—published their most operationally useful work yet, building a coherent set of guidelines for working with AI at scale.

Ruben Hassid's “27 Claude tips after 1,800 hours” was the week’s most shareable practitioner piece—a densely practical list built on genuine heavy usage going back to March 2023. Several arguments deserve unpacking. His first challenges the instinct to use **Claude Projects** for everything: Projects work well for repeated tasks with fixed material (client reports, standard procedures) but are actively counterproductive for creative work, because the model answers from your files rather than reasoning freshly. His tip on conversation length is counterintuitive but well-grounded: Claude re-reads the full conversation

before every response, so long conversations become expensive and the model begins averaging contradictions. He recommends starting a fresh chat after 30 to 50 turns, and discarding any thread over 100 turns entirely.

His most interesting formulation describes a skill as “a note you leave for a worker who may not ask a follow-up question before getting started”—meaning the editorial choices baked in by the skill’s creator arrive in your workflow whether or not they match your intentions.

“

A skill is a note you leave for a worker who may not ask a follow-up question before getting started.

— RUBEN HASSID

Nate's three newsletters this week built a coherent arc. His Sunday briefing used **Gumroad** as a case study: an agent found a charting bug, wrote a test, shipped the fix—then got the design wrong in a way only the human customer could catch, precisely because taste and context lived outside the agent's scope. His Wednesday token-saver piece surfaced a striking operational finding: 95.73% of his input tokens across 143 **Codex** threads were classified as “reused input.” His response: route to smaller models where appropriate, cache repeated input aggressively, cap context history, and set hard budget limits rather than soft alerts. His Saturday piece introduced a “one-job test” for evaluating any installed skill—seven steps from naming the job to a keep/fork/delete decision. His warning: past 25 installed skills, Claude Code and Codex begin silently trimming the skill list, and the model starts averaging across conflicting instructions. Adding skills to fix bad output is often what made the output bad in the first place.

Lenny Rachitsky's How I AI episode reviewed **Claude Opus 5** and demonstrated browser use inside **Codex** through five real examples, including an agent that used iPhone screen mirroring to update a router's firewall settings remotely while travelling. His interview with **Dianne Penn, Anthropic's** first technical PM and now Head of Product for AI Research and Labs, offered a rare inside view of the company's “eval-driven development loop”—where every capability improvement is anchored to measurable benchmarks before shipping. Penn joined in 2023 when the entire product team was five engineers. Nate's Monday model-evaluation piece offered a cautionary data point: in a 34-task run with full logs, one worker reported 213 verified quotations, and 13 turned out to be stitched together from fragments. A cheaper model can double review time, making total cost higher than the premium alternative even at one-fifth the token price.

SHOULDERS OF GIANTS

Creator Economy & the Self-Made Myth

Galloway's Board of Regents appointment becomes the launch pad for a systematic dismantling of Silicon Valley's favourite fiction—while 57% of Gen Z still want to be creators.

57%

GEN Z WHO WANT TO BE CREATORS

SOURCES: PROF G / SCOTT GALLOWAY · PROF G POD

The Infrastructure Nobody Built Alone

Galloway's "Shoulders of Giants" essay arrives at an unusual moment of honesty: the self-made myth is not merely morally wrong—it corrodes the infrastructure that makes future success possible.

Scott Galloway spent much of the week thinking about credit, capital, and who actually deserves either. His Monday "Young Money" post previewed a live discussion with **Jack Raines** on building wealth in your twenties. His Friday "Shoulders of Giants" essay was the more important work—and perhaps the week's most honest piece of writing from any source.

The essay opens with Galloway's appointment to the **University of**

California Board of Regents and an unusually direct accounting of unearned advantage: immigrant parents, a devoted single mother, "the unearned advantage of being born straight, white, and male during the largest economic expansion in history," and taxpayer-funded public education. He uses this personal frame as the launch pad for a systematic dismantling of the self-made myth as practised by Silicon Valley's most prominent voices.

“

The dishonesty is not just factually wrong—it corrodes the public investment in collective infrastructure that makes future entrepreneurs possible.

— SCOTT GALLOWAY, PROF G

Elon Musk claims zero deaths resulted from his DOGE-driven foreign aid cuts (against a conservative estimate of 750,000, two-thirds children), while having built **Tesla** on \$2.5 billion in US government support and relying on **NASA** for 20% of **SpaceX**'s revenue. **Chamath** told CNBC he would give away 95% of his wealth to taxes if government were effective—while having built his fortune through a company created on government-funded internet infrastructure. **Marc Andreessen** coded **Mosaic** at the publicly-funded University of Illinois while working at the federally-funded National Centre for Supercomputing Applications. Galloway's argument is structural, not merely moral: the dishonesty corrodes the public investment in collective infrastructure that makes the next generation of entrepreneurs possible.

The **Prof G Pod**'s Wednesday episode, "How the Creator Economy Took Over Media," picked up a complementary thread. 57% of **Gen Z** say they would like to be a creator for a living, but only a tiny fraction ever make real money at it. The top exists and is real: **Derek Thompson** passed his *Atlantic* salary five months after going independent; **Oliver Darcy** is closing in on 100,000 subscribers at *Stat*. But even those success stories rest on infrastructure nobody built alone: bandwidth, payment rails, platform reach, and decades of taxpayer-subsidised communications research. The creator economy took over media, Galloway argues—but it did so standing on the shoulders of giants who were, themselves, standing on the shoulders of the public.

About This Digest

PUBLICATION

The Elliot Edition is a private weekly digest edited for Elliot Gilford. It synthesises eight selected newsletters across technology, finance, geopolitics, and culture. Issue No. 2 covers the week of July 26 through August 2, 2026. Distribution is strictly private and not for redistribution.

SOURCES THIS ISSUE

Prof G Markets (Scott Galloway & Ed Elson) · DealBook (Andrew Ross Sorkin, Peter Coy) · Diplomatic (Laura Rozen) · GTMnow · Chamath Palihapitiya / All-In & Social Capital Research · GMO Asset Allocation · MastersInvest · Ruben Hassid · Nate (AI practitioner newsletters) · Lenny Rachitsky / How I AI · Prof G Pod · ARK Invest Q2 2026 Update.

TYPOGRAPHY

Display headings set in **Fraunces** (variable, SOFT axis), a wonky-serif designed by Undercase Type & Phaedra Charles. Editorial accents and pull-quotes in **Cormorant Garamond** (variable, Christian Thalmann). Section kickers and table of contents in **Archivo** (variable, Omnibus-Type). Body text in **Work Sans** (variable, Wei Huang). Labels, folios, and metadata in **Space Mono** (Colophon Foundry). Poster kickers in **Anton** (Vernon Adams). All fonts embedded under their respective open-source licences.

PRODUCTION

Typeset in HTML & CSS and rendered to print-ready PDF via WeasyPrint 69. Saddle-stitch booklet imposition for A4 duplex printing via pypdf 6. All layout and editorial work performed by Claude (Anthropic), August 2026. No external imagery or illustration. Colour system: Cream #f2ead9, Ink #241f18, Burgundy #6d2438, Navy #263449, Rose #c48a72.

PRIVATE DIGEST · NOT FOR DISTRIBUTION

The Elliot Edition

*A weekly synthesis of what matters across
technology, markets, geopolitics, and culture.*

ISSUE NO. 2 · JULY 26 – AUGUST 2, 2026



ELLIOT-ED-002-2026