

Weekly Newsletter Summary: July 26 – August 2, 2026

This was a week that put the entire AI investment thesis under stress — chips melted down, a high-flying hedge fund nearly collapsed, the Iran war kept grinding away at bond markets and polling numbers, and yet the practitioners building with AI quietly published some of the most useful operational guidance yet. Across eight themes, here is what your newsletters were thinking about.

Big Tech's Hidden AI Debt and the Chips Meltdown

The week opened with what Scott Galloway and his markets co-writer Ed Elson called the most alarming finding they had covered in months: an investigation by Nikkei Asia revealing that Alphabet, Amazon, Meta, Microsoft, and Oracle are carrying one point six five trillion dollars in off-balance-sheet debt that investors are largely not counting. The mechanism is relatively straightforward but its scale is staggering. Instead of building data centers themselves, the hyperscalers sign long-term lease agreements with data center operators — operators that are often shell companies funded by private credit providers like Blue Owl Capital, Apollo, and Blackstone. Those private credit funds are in turn financed by pension funds and life insurance companies. Galloway points out that the California State Teachers' Retirement System is one of Blue Owl's biggest investors. The implication is that if the AI bubble bursts, the risk does not fall on Big Tech's balance sheets. It falls on the retirement savings and insurance policies of ordinary working people, while staying invisible to the SEC and Federal Reserve, who have no regulatory jurisdiction over private credit.

DealBook picked up a parallel thread on Tuesday with its "Chips meltdown" edition. South Korea's Kospi index fell more than ten percent in a single session, wiping nearly six hundred billion dollars from SK Hynix's market cap. Nvidia tumbled five percent in the US, losing its crown as the world's most valuable listed company — to Apple. Andrew Ross Sorkin identified three forces at work: anxiety about whether the AI capex binge is generating real returns, concerns about circular financing (Nvidia is reportedly in talks to provide a two-hundred-and-fifty-billion-dollar guarantee for OpenAI's Ohio data center, raising questions about whether chipmakers are essentially vendor-financing their own biggest customers), and growing Chinese competition. Goldman Sachs estimates the four biggest hyperscalers will spend a combined five point three trillion dollars on

capital expenditure through 2030, and they plan to borrow up to four hundred billion dollars this year alone to do it.

By Wednesday, when DealBook ran its “Big Tech turmoil” edition, the picture had become more complicated. Microsoft reported a quarterly profit jump of thirty-one percent to thirty-five point eight billion dollars, with cloud revenue growing at its fastest rate in four years. Meta, by contrast, fell sharply in premarket trading. The contrast illustrated the central investor anxiety: Microsoft is already generating significant AI revenue from its cloud unit; Meta is still spending heavily without the same evidence that it is paying off.

The sharpest analytical framing of the whole chips story came from GTMnow’s Friday piece on the token price collapse. Their core argument deserves to be listened to carefully. The price per token has fallen more than ninety-five percent in three years — OpenAI’s flagship launched at thirty dollars per million input tokens in March 2023, and today a comparable task costs under fifty cents. And yet enterprise spend on large language models more than doubled in six months, from three point five billion dollars in late 2024 to eight point four billion by mid-2025. Sundar Pichai revealed at Google I/O 2026 that Google went from processing nine point seven trillion tokens per month to over three point two quadrillion — a three-hundred-and-thirty-fold jump — over the same period that per-token prices were falling. GTMnow’s framing draws on the Jevons Paradox: when a better steam engine cut the coal needed per job, total coal use went up. Cheap tokens did the same thing. The article goes further, arguing that this is why AI-native businesses structurally run at fifty to sixty percent gross margins rather than the eighty to ninety percent typical of traditional SaaS — every inference call is a variable cost of goods sold. One cautionary tale: a four-agent system that had no hard cost cap looped for eleven days and burned forty-seven thousand dollars before anyone caught it.

Then on Friday morning came the most dramatic story of the week. DealBook reported that Situational Awareness, the hedge fund founded two years ago by Leopold Aschenbrenner — a former OpenAI researcher with no prior investment experience — had seen its portfolio of mostly AI bets fall roughly sixty-seven percent in a single month, forcing a fire sale to Ken Griffin’s Citadel to cover margin calls. Aschenbrenner had raised the fund on the strength of his 2024 essay predicting AI would transform society, attracting backing from Stripe founders Patrick and John Collison and former GitHub CEO Nat Friedman. The fund raised over ten billion dollars, held a three point five billion dollar stake in Anthropic, and had been up around eighty percent on the year before this month’s collapse. Roger Lowenstein, author of “When Genius Failed,” told DealBook the situation felt “a little more like the dot-coms — huge equity investments in a new thing that no one can value with any hope of precision.” The Citadel lifeline appeared to contain the immediate damage, with AI stocks rebounding on Friday — but the question of how much hidden leverage

exists in the AI trade, which Galloway raised earlier in the week through his hidden-debt piece, suddenly felt very immediate.

The Iran War Drags On: Politics and the Bond Market

The geopolitical backdrop running through every market conversation this week was the ongoing US-Iran conflict, now approaching six months since Trump's announcement. Laura Rozen's Diplomatic newsletter delivered the clearest political read on Wednesday: Trump has hit an all-time high disapproval rating of sixty-six percent in a new CNN poll, with only thirty-four percent approval. The Quinnipiac poll found thirty-two percent approval, fifty-eight percent disapproval — Trump's lowest Quinnipiac number ever. More significantly, Republican approval has eroded from eighty-five percent in June to seventy-six percent now. Democratic consultant David Doak told Rozen that what is unusual about this moment is not just that voters are unhappy, but that they have specifically blamed Trump for causing the problems, concluded he lacks a strategy in Iran, and further concluded — most damaging of all — that “he doesn't care.” Doak argues that these three things — assignment of blame, diagnosis of incompetence, and the sense of indifference — are the hardest combination to recover from politically.

DealBook's Thursday “Credibility shock” piece connected the war directly to the bond market. The thirty-year Treasury yield rose to five point two three percent — the highest since 2007, right before the global financial crisis. Traders were pricing in sixty percent odds of a September interest rate increase. Kevin Warsh, the Fed chairman, held rates steady in a dissent-filled nine-to-three vote but sent mixed signals: he opened the door to measuring inflation against indicators beyond PCE, suggested markets had already done some tightening for the Fed, and simultaneously insisted the Fed would not go “soft” on its two percent target. Bank of America economist Aditya Bhave wrote that the market moves were “consistent with a central bank inflation credibility shock.” Scott Galloway's Monday markets piece put numbers on the broader fiscal vulnerability: every one-tenth of a percentage point rise in rates over the coming decade increases US interest costs by roughly one point eight trillion dollars, against a national debt already at thirty-nine point four trillion dollars. Interest payments on that debt already cost nearly one trillion dollars in fiscal year 2025 — about seven thousand seven hundred dollars per American household — and are projected to become the second-largest federal expenditure this year.

The AI Open vs Closed Debate

If there was a single ideological battle dominating the tech conversation this week, it was the open versus closed AI debate, which exploded into full public view on Monday when Jensen Huang posted on X to announce the Open Secure AI Alliance — a coalition including Andreessen Horowitz, Meta, and Microsoft arguing that open models would accelerate innovation and strengthen security. The post attracted more than sixty million views. Huang’s framing was deliberately inclusive: “The world needs both frontier closed models and frontier open models.” But the subtext was unmistakable. As DealBook noted, Nvidia has financial skin in the closed camp — it is a major backer of both Anthropic and OpenAI — which makes Huang’s public embrace of open models a strategically interesting move. His logic, analysts suggested, is that easier access to open-source AI speeds the technology’s growth and drives demand for Nvidia’s chips regardless of which labs win.

Mark Zuckerberg was blunter. He told the Times that Anthropic and OpenAI’s argument for keeping powerful models in controlled environments was “bad for innovation and even dangerous,” adding that “so much of the discourse from a lot of the other labs that are developing this is overwhelmingly filled with doom.” Chamath Palihapitiya, writing his Sunday newsletter in response to the distillation controversy, took the same side but with a more cynical framing. The labs claim that Chinese competitors like DeepSeek and Moonshot AI are stealing their technology through distillation — running their models at scale to harvest question-and-answer pairs for training. Chamath’s counter-argument is that if this were a genuine security concern rather than a competitive pretext, the labs would solve it with KYC — know-your-customer verification, real identities, bounded credit cards — which would kill industrial-scale account farms overnight. They do not do this, Chamath argues, because it would slow revenue. His conclusion: “What you’re watching is valuation preservation.” He goes further, arguing that the real moat is not in the model layer, which is commoditizing faster than anyone expected, but above it (applications that people pay for) and below it (chips, cloud, infrastructure). DealBook noted that Anthropic and OpenAI accused Chinese start-ups of IP theft via distillation, while Treasury Secretary Scott Bessent threatened sanctions against any Chinese companies that “cross the line.”

Autonomous Vehicles: The Coming Revolution

Chamath Palihapitiya’s Friday deep dive on autonomous driving was the week’s most ambitious single piece — a ninety-two-page research report from Social Capital summarizing where the technology actually stands. The numbers he leads with set the stakes: the average American worker

spends roughly two hundred and twenty-five hours per year commuting. In 2025, there were over thirty-six thousand traffic fatalities on US roads. Waymo's peer-reviewed safety data across one hundred and seventy million driverless miles shows ninety-two percent fewer serious crashes or fatal injuries than human drivers in equivalent conditions. Waymo passed twenty million cumulative paid driverless rides by end of 2025 and is now running five hundred thousand rides every week. Autonomous ride-hailing is commercially available in twenty-eight cities across five countries.

The economic thesis hinges on one metric Chamath calls "the price of a mile." A personally owned vehicle costs approximately one dollar and ten cents per mile when you account for insurance, maintenance, depreciation, and fuel. ARK Invest estimates the robotaxi market at eleven trillion dollars once cost per mile hits twenty-five cents. Goldman Sachs expects sixty-five percent of US car sales could be autonomous vehicles by 2040. Aurora plans two hundred or more driverless trucks by the end of 2026, targeting the nine-hundred-billion-dollar US trucking industry where labor accounts for forty-five percent of per-mile costs. Chamath's report also highlights the second-order industry disruptions: car insurance, personal injury law, and traffic enforcement all face structural rethink when human error is largely removed from the equation.

Google's Record Quarter and the Space Economy

Chamath's Sunday piece, published under the headline "Google's biggest quarter ever," was actually more focused on the open-source distillation debate than on Google's financials — the Alphabet earnings context served as the launch pad for his broader argument about AI commoditization. But the business backdrop was significant: Google's token processing volume, as Sundar Pichai revealed at I/O 2026, had grown three-hundred-and-thirty-fold in two years, even as per-token prices fell continuously. That is the cleanest single data point illustrating the Jevons Paradox dynamic that GTMnow would elaborate on the following Friday.

The space economy got its own dedicated treatment in DealBook's Saturday "Space grab" piece, written by Peter Coy. The central story is that SpaceX is racing to acquire radio spectrum — the invisible real estate that carries wireless signals — in a contest that Coy compares to the land rushes of the nineteenth century. SpaceX's IPO prospectus, which valued the company at one point eight trillion dollars, flags spectrum acquisition as a major risk factor, noting that some countries are considering changing how they grant domestic operating licenses in ways that would favor homegrown operators over American companies. China has been flooding the International Telecommunication Union with spectrum applications while simultaneously advocating for

“equitable access” — a defensive move against SpaceX’s first-mover position. The coming battleground is the World Radiocommunication Conference in Shanghai in October and November 2027. Meanwhile, the ARK Venture Fund announced an investment in K2 Space’s Series D — a vertically integrated satellite manufacturer building high-power spacecraft for commercial and national security missions — and its Q2 2026 update noted that the successful public listings of SpaceX and X-energy marked an “inflection point” for the IPO market reopening to category-leading innovation companies.

Markets and Investing

The week’s market-oriented newsletters offered three complementary angles on portfolio strategy under unusual stress. GMO published its latest seven-year asset class forecast on Monday, a regular update to its well-known long-horizon return projections, without elaborating detail in the newsletter itself. More substantive was GMO’s Wednesday piece on what they called “the liquid alternatives revival” — a direct challenge to how most investors think about diversification. The argument runs as follows: investors have historically relied on bonds to diversify equity risk and on illiquid alternatives (private equity, hedge funds, etc.) for return potential. GMO’s asset allocation team argues that bonds no longer reliably diversify in every scenario — particularly in an inflationary, wartime environment like the current one — and that the macro conditions supporting private equity’s outperformance have ended. Illiquid alternatives, they contend, have been failing to live up to their inflated promises. Liquid alternatives — hedge fund-like strategies accessible in a liquid fund format — may offer a compelling combination of diversification, liquidity, and return potential that the traditional sixty-forty portfolio cannot.

The MastersInvest “Investment Masters Class” contribution this week came under the evocative title “Waiter or Owner,” drawn from an AB InBev executive’s observation that most companies inadvertently create employees who think like waiters rather than like restaurant owners. The piece traces employee ownership models through UPS (where Jim Casey first offered employee shares in 1927, calling it “the single smartest move”), Publix (which has annual voluntary turnover of around five percent against an industry average of sixty-five percent, largely because store associates share directly in store profits every thirteen weeks), Walmart, and HEICO. The investment thesis is that genuine ownership alignment — where employees receive meaningful economic stakes and the authority to act on them, not just motivational language about ownership culture — produces businesses that sustainably outperform because the people closest to the customer have an actual financial stake in what happens.

AI Tools, Claude, and the Practitioner Edge

This section of the week's reading was where you could hear the gap widening between people who are extracting real leverage from AI tools and those who are still treating them like search engines with better grammar. Ruben Hassid's "27 Claude tips after 1,800 hours" piece was the most widely shareable of the bunch — a densely practical list built on genuine heavy usage that started back in March 2023. A few of his arguments deserve to be unpacked rather than just listed. His first tip challenges the instinct to load everything into Claude Projects: he argues that Projects are great for repeated tasks with fixed material (client reports, standard procedures) but actively bad for creative work, because the model answers from your files rather than thinking. His tip on conversation length is counterintuitive but well-reasoned: Claude re-reads the entire conversation before every response, so long conversations become expensive and the model starts averaging out contradictions. He recommends starting a fresh chat after thirty to fifty turns, and leaving anything over a hundred turns behind entirely. His most interesting framing is the analogy of a skill to "a note you leave for a worker who may not ask a follow-up question before getting started" — meaning the tasteful decisions baked into a skill by its creator arrive in your workflow whether or not they match your definition of good output.

Nate's three newsletters this week built a coherent arc around the same fundamental problem. His Sunday executive briefing used Gumroad as a case study: a customer support agent found a charting bug, wrote a test, shipped the fix — and then got the design wrong in a way only the human customer could catch. His point is not that the agent failed but that it succeeded at everything within its scope and hit a natural boundary at the moment when taste and context belonged to a human. His Wednesday token-saver piece addressed a different operational problem: Nate found that ninety-five point seven three percent of his input tokens were classified as "reused input" across one hundred and forty-three Codex threads — not necessarily waste, since continuity is genuinely valuable, but a signal worth investigating. His skill to address this enforces four changes: routing to smaller models where appropriate, caching repeated input aggressively, capping context history to prevent it from growing quadratically, and setting hard budget limits rather than alerts. His Saturday piece then zoomed out to the meta-question of how to evaluate any installed skill using what he calls the "one-job test" — seven steps from naming the job to rerunning it, ending in a keep/fork/delete decision. The key warning he delivers: past twenty-five installed skills, Claude Code and Codex begin trimming the skill list silently, and the model starts averaging across conflicting instructions. Adding skills to fix bad output is often the thing that made the output bad in the first place.

Lenny Rachitsky's How I AI episode this week reviewed Claude Opus 5 and demonstrated browser use inside Codex through five real examples — including a particularly striking one where an agent used iPhone screen mirroring to update a router's firewall settings remotely while the user was travelling. His interview with Dianne Penn, Anthropic's first technical PM (now Head of Product for AI Research and Labs), offered a rare inside view of how the company thinks about product development. Penn joined in 2023 when the entire product team was five engineers. She described the “eval-driven development loop” her team has pioneered — where every model capability improvement is anchored to measurable benchmarks before being shipped. She also discussed “token maxing” — the practice of throwing large amounts of compute at problems — and where human judgment remains irreplaceable even as model capability improves. Her framing of Claude's willingness to push back as “key to its success” is interesting in context: it suggests Anthropic has made a deliberate bet that honest, calibrated responses that challenge users will build more durable trust than ones optimised purely for user satisfaction in the moment.

Nate's Monday model-evaluation piece rounded out the AI practitioner section with a data-driven look at Chinese models. His framing is careful: he uses DeepSeek, Qwen, GLM, Kimi, and MiniMax selectively and aggressively in some workflows — but he insists that job, model, and deployment path are separate decisions. His bakeoff guide includes a validator, a manifest, a score sheet, and “two fixtures you use to prove your checker actually rejects bad work.” His most useful cautionary data point: in a 34-task run where he has full logs, one worker reported two hundred and thirteen verified quotations, and thirteen of them turned out to be stitched together from fragments. The lesson is that a cheaper model can actually double your review time, making total cost higher than the premium alternative even at one-fifth the token price.

Creator Economy and Media

Scott Galloway spent the week thinking about success, money, and who actually deserves credit for either. His Monday “Young Money” post previewed a livestream with Jack Raines on how to spend your twenties — building wealth, finding purpose, and navigating financial decisions without putting life on hold. His Friday “Shoulders of Giants” essay was the deeper work. It opened with Galloway's appointment to the University of California Board of Regents and an unusually direct acknowledgment of how much his own success traces to things he did not earn: immigrant parents, an unusually devoted single mother, the “unearned advantage of being born straight, white, and male during the largest economic expansion in history,” and taxpayer-funded public education. He used this personal reflection as the launch pad for a systematic dismantling of the self-made myth as

practiced by Silicon Valley's most prominent voices. Musk, who claims zero deaths resulted from his DOGE-driven foreign aid cuts (against a conservative estimate of seven hundred and fifty thousand, two-thirds children), built Tesla on two point five billion dollars in US government support, while NASA accounts for twenty percent of SpaceX's revenue. Chamath told CNBC he'd give away ninety-five percent of his wealth to taxes if government were effective — ignoring that the government-funded internet research created the conditions for the company through which he built his fortune. Andreessen coded Mosaic at the publicly-funded University of Illinois while working at the federally-funded National Center for Supercomputing Applications. Galloway's argument is that the dishonesty is not just factually wrong — it corrodes the public investment in collective infrastructure that makes future entrepreneurs possible.

The Prof G Pod's Wednesday "How the Creator Economy Took Over Media" piece picked up a complementary thread. Fifty-seven percent of Gen Z say they'd like to be a creator for a living, but only a tiny fraction ever make real money at it. The top of the market is real — Derek Thompson passed his Atlantic salary five months after leaving to write independently, Oliver Darcy is closing in on a hundred thousand subscribers at Stat — but Galloway's broader point from the week is that even those success stories rest on infrastructure that wasn't built by any individual. Lenny Rachitsky's "So, you want to be a content creator?" piece this week addressed the same question from a more practical angle, though the full content of the article was not available in the email preview.

Week in One Thought

The deeper pattern underneath everything this week — the hidden debt, the chips rout, the fund blowup, the open-vs-closed battle, the bond market stress, the practitioner scramble to tame costs — is that we are in the part of every great technological transformation where the initial enthusiasm has met actual economics, and the gap between what people believed was happening and what the numbers show is now being priced. The people who will come out ahead are the ones who can see both sides simultaneously: who understand that the technology is genuinely revolutionary while also recognising that the current financial structures built around it are fragile, overleveraged, and in some cases deliberately opaque about who bears the risk when the bet goes wrong.