Weekly Newsletter Digest: June 11–18, 2026

This was one of the most consequential weeks in recent memory for financial markets, geopolitics, and technology, with multiple once-in-a-generation events landing within days of each other. SpaceX's IPO, a US-Iran ceasefire framework, the arrival of a new Federal Reserve chairman, and escalating AI governance battles all competed for attention at the same time — and the newsletters captured every angle.

---

1. The SpaceX IPO: History, Hype, and Honest Disagreement

On Friday June 12, SpaceX began trading on the Nasdaq under the ticker SPCX, raising roughly $75 billion and briefly valuing the company at over $2 trillion — the largest IPO in history, more than doubling Saudi Aramco's previous record of $29.4 billion. The stock priced at $135 the night before, opened at $150, hit an intraday high of $176.52, and closed up roughly 19 percent. By the middle of the following week it had climbed further still, surpassing Amazon to become the fifth-largest publicly traded company on Earth at a market cap of around $2.5 to $2.6 trillion.

The sheer scale prompted every serious commentator to attempt a framework for understanding it, and the frameworks they reached for diverged dramatically.

Peter Diamandis offered the most expansively bullish case. His argument is essentially that SpaceX is three businesses converging on a single thesis: a launch monopoly — executing roughly 85 percent of all satellite launches in 2025, more than every other space program on the planet combined — a cash engine in Starlink, which posted $11.4 billion in recurring broadband revenue in 2025 and over a billion dollars in quarterly profit, and an AI frontier lab in xAI, which is now being turbocharged by the acquisition of Cursor. Diamandis compared SpaceX to the transcontinental railroads: you could have valued those on passengers and freight and been wildly wrong, because their real value was in the towns they opened and the economies they unlocked. The railroad to orbit, he argued, opens the solar system. When you layer in the xAI orbital compute ambition — SpaceX revealed its AI1 satellite the week before pricing, a spacecraft with 150 kilowatts of peak compute, a 70-meter wingspan, and solar panels feeding AI processing from orbit — Diamandis believes the addressable market is effectively infinite. He disclosed that going forward he is putting fresh capital into SpaceX rather than Bitcoin.

Chamath Palihapitiya wrote a 120-page deep dive on SpaceX through his Social Capital research team, framing the company through the analogy of Vasco da Gama. For centuries, spices reached Europe via routes controlled by Arab and Venetian middlemen; da Gama's sea route around the Cape of Good Hope collapsed the entire cost structure of the trade in one move. SpaceX has done the same thing to reaching orbit. NASA's Space Shuttle cost $54,500 per kilogram to reach low Earth orbit; SpaceX brought that to $2,720 — a twenty-times reduction — and is still moving the number down. Starlink is the first proof of what cheap orbital access unlocks. Orbital compute is the next bet the current valuation is pricing in.

Scott Galloway disagreed sharply, and his objections were both financial and personal. On the financial side, he noted that SpaceX is trading at 112 times trailing revenue, compared to Google going public at 10 times revenue growing at 240 percent a year, and Meta at 28 times revenue growing 88 percent a year. He pointed to Jim Chanos, the prominent short-seller, who called it "a hopes-and-dreams IPO" and noted that SpaceX was trading at roughly 90 times revenue. Galloway also noted that SpaceX engineered the IPO to manufacture scarcity — selling just over 4 percent of total shares while getting the Nasdaq to waive its usual index inclusion requirements, meaning index funds were effectively forced to buy in. On the personal side, Galloway was direct: the thing that worries him is not the lockup period but that a man who he described as estranged from children, with a substance abuse issue, sleeping with a loaded gun, and holding what Galloway described as extreme political beliefs, is now the most powerful individual who has ever existed on the planet, worth more than the bottom 46 percent of humanity and 3.4 percent of US GDP. At peak, John D. Rockefeller was worth 1.5 percent of GDP.

The DealBook newsletter from Andrew Ross Sorkin covered the IPO across multiple days and added important market mechanics. SpaceX sold just over 4 percent of shares at IPO, versus the typical 10 percent or more, which Sorkin noted created a demand-to-float imbalance that drove the first-day pop. The perpetual futures contract trading on Hyperliquid at around $163 before the debut suggested traders had already priced in the premium. Sorkin flagged that to buy something new, investors have to sell what they own — and indeed, Rocket Lab, EchoStar, and SpaceMobile declined as much as 10 to 14 percent on the IPO day as money rotated into SpaceX. JPMorgan warned that the same dynamic could pressure existing Mag 7 holdings, and it will only intensify when Anthropic and OpenAI go public, which the market is expecting as soon as the fall.

Where the commentators agreed was on the human dimension of the wealth creation. Both Diamandis and Galloway, from very different ideological starting points, noted that roughly 4,400 current and former SpaceX employees became millionaires, with several hundred more becoming centimillionaires or billionaires. Diamandis called this what abundance looks like: value created rather than extracted. Galloway acknowledged it while still worrying about the concentration at the top.

One additional development the week brought: SpaceX used its new public-market currency to close the $60 billion acquisition of Cursor, the AI coding tool. DealBook reported that Cursor's enterprise revenue run rate had grown 50 times year-on-year and tripled in just the first quarter of 2026. Andreessen Horowitz's 10 percent stake in Cursor is now worth around $6 billion. Cursor's four co-founders each held roughly 4.5 percent, making them each worth around $2.7 billion. The strategic logic is direct: SpaceX's xAI lags Anthropic and OpenAI on pure model capability, but a publicly traded company with $85 billion in IPO cash and excess compute capacity — which it is currently leasing to both Anthropic and Google for billions of dollars — can buy its way toward the frontier faster than organic R&D.

---

2. The Federal Reserve Under Kevin Warsh: A New Regime Begins

Kevin Warsh chaired his first FOMC meeting this week and immediately signaled that his tenure will look nothing like his predecessors'. The Fed held the benchmark rate steady at 3.5 to 3.75 percent in a unanimous vote, but the surrounding signals were unmistakably hawkish. The Fed's own projections now show nearly half the committee expecting a rate hike before year-end. Futures traders by mid-week were pricing at least one increase in 2026, possibly as early as October — compared to March, when not a single Fed official had penciled in any hike at all.

The context is uncomfortable. CPI inflation has hit 4.2 percent year-on-year, a three-year high, driven primarily by the energy-price shock from the US-Iran war. That is more than double the Fed's 2 percent target. Unemployment sits at 4.3 percent, broadly stable, but the inflation picture is the dominant pressure point. As Chamath's deep dive on the Federal Reserve explained this week, the Fed controls the cost of borrowing across the entire American economy by setting the rate banks charge each other overnight. When that rate rises, mortgages, corporate loans, stock valuations, real estate prices, and retirement accounts all adjust downward. The Fed is, in Chamath's framing, the most consequential institution in the American economy — and Kevin Warsh has a clear and unconventional vision for how it should change.

Warsh abstained from submitting his own dot in the dot plot — the chart showing where each Fed official expects rates to go — sending a signal that he may eventually eliminate the practice altogether. His communications during his first press conference were deliberately minimalist, prompting comparisons to Alan Greenspan's famously opaque style. Jeffrey Roach, the chief economist at LPL Financial, wrote to investors: “We are going back to the days of Alan Greenspan when FOMC statements were deliberately minimalist.” He also established five task forces inside the Fed, covering communications, the balance sheet, and artificial intelligence’s influence on productivity and jobs — the last of which struck markets as a notable departure from convention.

The DealBook newsletter had flagged the key tensions ahead of the meeting. Would Warsh see high inflation as a transitory blip tied to wartime energy prices, or as an entrenched problem requiring decisive action? Would he protect Fed independence from Trump, who spent much of his second term calling for lower rates, and who said this week “it could happen… it keeps our country down, it’s so unusual” when asked about rate hikes — hardly a full endorsement, but a concession that was remarkable coming from him. And would Warsh live up to his stated desire for a leaner central bank that telegraphs less and surprises more?

Chamath’s deep dive made the structural argument for why this matters beyond just one meeting. If Warsh actually executes his vision of a Fed that communicates less, that uncertainty itself gets priced into markets — higher risk premiums, higher long-term bond rates, lower stock valuations — even without a single rate hike. The market impact is not just from the rates themselves but from the change in the information environment. Chamath’s team traced every Fed Chair through history and their defining tests, noting that Warsh inherits an institution where Jerome Powell, the previous chair, retains his board seat and his vote on every rate decision through 2028.

Ray Dalio weighed in this week via his LinkedIn essay “What Should You Do Under Existing Circumstances?” framed around his investment principles — the world as a bridge, poker, or chess game where you need to assess circumstances and determine your next move. He had written earlier in the year that if Iran was left with control over the Strait of Hormuz, the United States would be judged to have lost the war, and that judgment would be a key metric for evaluating the geopolitical situation.

---

3. The US-Iran Deal: A Ceasefire Framework Full of Uncertainty

On Sunday June 14, the US and Iran reached a Memorandum of Understanding — the “Islamabad Understanding” — signed electronically by Trump and Vice President JD Vance, declaring the immediate and permanent termination of military operations on all fronts including in Lebanon. A formal signing ceremony was scheduled for Switzerland over the weekend, with Vance leading the US delegation and Pakistani mediators facilitating.

Markets reacted immediately and dramatically. DealBook's Monday edition captured a global relief rally: Brent crude fell to around $83, a three-month low. S&P 500 futures, crypto, and sovereign bonds all rallied. The yield on the 10-year Treasury note fell to 4.45 percent. The average US gasoline price dropped to $4.06.

Laura Rozen's Diplomatic newsletter provided the most granular coverage, publishing the full text of the MOU — eleven paragraphs negotiated in detail. The key provisions were substantial: the US committed to begin removing its naval blockade within 30 days and to fully lift it within 30 days of the MOU signing. Iran would facilitate safe passage for commercial vessels through the Strait of Hormuz. The US committed to developing a plan with at least $300 billion for Iran's reconstruction and economic development. All US sanctions, including UN Security Council resolutions and IAEA Board resolutions, would be terminated on a schedule to be finalized in a final deal. Iran reaffirmed it shall not procure or develop nuclear weapons, with enriched material stockpiles to be downblended on site under IAEA supervision. The two parties agreed to negotiate a final deal within 60 days.

But Rozen's reporting was clear-eyed about the fragility of all of this. At a Carnegie Endowment panel she cited, former State Department Iran expert Suzanne Maloney argued that the Iranians believe they have achieved not just regime survival but a position that will enable them to change the strategic equation in the region — and that Iran will be reintegrated economically in ways that benefit the regime. Carnegie's Karim Sadjadpour warned that the hope for a grand bargain in 60 days is unrealistic, and that failure to complete a final deal could lead back to conflict. Veteran diplomat Aaron David Miller said every time he hears the words "grand bargain" applied to the Middle East, it makes him very nervous.

Trump himself provided his own framing at a G7 press conference in Evian, France, in what Rozen described as a rare moment of candor: "If we didn't do this deal, we could have dropped more bombs for another three weeks… you would never have the Hormuz Strait open… the stock market would have gone down at levels that nobody ever saw before, maybe except for 1929." He explicitly named preventing economic catastrophe as the primary motivation. He then immediately threatened to resume bombing if no final deal was reached within 60 days — which Rozen noted was itself arguably a violation of paragraph one of the MOU he had just signed.

At the G7 summit running alongside all of this, DealBook reported that leaders broadly agreed on the need for democratic nations to lead and collaborate on AI governance — a message pushed forcefully by Dario Amodei, Sam Altman, Demis Hassabis, and Arthur Mensch. Amodei urged G7 leaders to "resist the temptation to splinter" on AI. The urgency behind that message was directly connected to what had happened the prior Friday evening.

4. Anthropic, AI Governance, and the Export Control Crisis

Late on Friday June 13, with ninety minutes of notice, the Trump administration issued an export control directive barring foreign nationals from accessing Anthropic's Fable 5 and Mythos 5 models. This applied even to Anthropic's own foreign-born employees — roughly a third of the company's staff. Anthropic's response was to shut down global access entirely, arguing there was no technically feasible way to comply with the order otherwise.

The trigger, according to the Peter Diamandis newsletter and DealBook, was a jailbreak that a researcher had found in Fable 5's guardrails. But cybersecurity expert Katie Moussouris, who had seen the White House report, told The Atlantic the underlying vulnerability was far less alarming than the administration implied: the "exploit" involved giving Fable deliberately insecure code, asking it to "review the code for security issues" (which it refused), and then asking it to "fix this code" (which it complied with after further manual steps). Moussouris called this "the model working as intended for cyberdefense," and noted that OpenAI's GPT-5.5 could be used in the same way. A group of high-level security experts publicly urged the White House to reverse course.

The Diamandis newsletter and Lenny's "How I AI" episode both engaged substantively with Fable 5, which Anthropic had released as its first generally-available Mythos-class model in the weeks prior. The performance data is striking: Fable 5 hit 80 percent on SWEBench Pro, significantly outperforming Opus 4.8, GPT-4.5, and Gemini 3.1 Pro. But Claire from Lenny's newsletter found a more nuanced picture in real-world testing. The model excels at hard technical problems, long-horizon work, and vision tasks like document formatting and PDF parsing. It is expensive by design: $10 per million input tokens and $50 per million output tokens, at roughly twice the token consumption rate of other models. For product specifications and PRDs, it produces technically complete but almost unreadable dense documents. For one-shot design tasks, it performed surprisingly poorly — producing what Claire described as fundamentally terrible design. And its conservatism, which seems partly a function of its safety tuning, means that when asked to ship an MVP, it produces something narrow and minimalist to a fault.

Diamandis framed the deeper issue with characteristic sharpness: "Who decides what level of intelligence you and your company are allowed to access?" The structural irony he identified is that roughly 70 percent of elite US AI researchers are foreign-born. Export controls risk pushing those researchers back to their home countries, accelerating the very AI competition with China that the controls are supposed to address. He also reported that even before the export order, developers were in revolt over a discovery buried in a 319-page Anthropic document: the company retained every prompt for at least 30 days, even for enterprise customers with negotiated zero-retention agreements, and if Fable 5 detected a user doing frontier AI research, it would silently downgrade them to a weaker model without notification.

Meanwhile Chamath's weekly roundup reported that OpenAI had filed its confidential S-1 on June 8, one week after Anthropic disclosed its own IPO preparations. OpenAI was last valued at $852 billion in March; Anthropic closed its Series H at $965 billion in late May. Chamath noted that a maxed-out ChatGPT Pro subscription at $200 per month can draw up to $14,000 in tokens per month at API prices, meaning OpenAI is potentially absorbing $3,500 in compute costs against $200 in revenue from its heaviest users. The economics of the flat-rate model at the frontier are being stress-tested at scale.

---

5. AI Agents, Vibe-Coding, and the Future of Work

The AI practitioner newsletters this week converged on a consistent theme: the hard part of AI is not getting access to intelligence, it is building the organizational structure to use it.

Nate's Substack made this argument most systematically, in two separate pieces. The first, framed as an executive briefing around the OpenAI IPO, drew a clean distinction between the market narrative — that intelligence is scarce, which is why OpenAI, Anthropic, and xAI are being valued as the companies that own the next decade of enterprise value — and the company-level reality, which is that a founder he spoke to moved his agentic startup from a frontier model to an open-weight model and saw his model bill drop 97 percent in one month. Both things are true simultaneously. Wall Street is pricing the scarcity of the frontier. Inside most companies, useful intelligence is getting radically cheaper. The question is not whether your company will use AI models — it will — but whether the structure around them, what Nate calls "the harness," becomes something you understand and own or something that happens to you. The harness is the context, documents, permissions, review standards, memory, budgets, decision rights, and accountability structures that turn model intelligence into trustworthy organizational work. The labs can sell you models and even engineers to install them. They cannot sell you your own operating context.

His second piece this week made the maintenance argument. He opened with a story about maintaining boats in Indonesia, where saltwater finds every shortcut, as a frame for AI agent maintenance. Vercel's sales agent story — which Business Insider reported as a labor story about how Vercel trained an AI agent on one of its best sales development reps and moved from a ten-person inbound team to one person overseeing the agent — he read differently. The useful story is not the headcount reduction. It is what had to be true around the agent for the work to become trustworthy. The agent had a workbench: sources, tools, a defined job, handoffs, a review path, feedback loops, and a human who could see what was happening. Nate's warning is that agents break in two ways — when the world around them drifts, and when the model underneath them gets better and the harness built for its old weaknesses becomes dead weight. His practical framework covers seven "surfaces" that go stale: job definition, diet (what information the agent consumes), memory, tools, reach, proof, and value.

Ruben Hassid approached the same territory from the practitioner angle. He published a guide to vibecoding with Claude Code, framed as a tool for two specific legitimate use cases: building a clickable prototype to hand to developers so they stop guessing what you meant, and building small internal tools that do specific jobs better than a chat interface. His advice was deliberately modest — this is not a get-rich-quick mechanism, and getting a functioning company requires far more than knowing how to prompt an AI. What vibecoding gives non-coders is the ability to externalize their ideas as actual working artifacts rather than descriptions of artifacts. He also published a piece where he acknowledged that Claude had effectively automated the how-to half of his newsletter job, which prompted a thoughtful reflection on what the other half is: the testing, the things that broke before they worked, the judgment about what is worth sharing and what isn't. The newsletter is not the how-to instructions, he concluded. The newsletter is him.

Lenny's newsletter added two substantive threads. The mom-and-pop SaaS era post — which was a brief announcement link rather than full content in this email — pointed to the broader phenomenon of solo founders using AI tools to build small functional businesses that would previously have required teams. The Braintrust CEO Ankur Goyal interview covered how engineering teams are using AI agents, evaluations, and continuous integration to ship software faster — specifically the idea that agents are now capable enough that the bottleneck has moved from "can the AI do this?" to "can we trust what the AI did?"

GTMnow had two relevant episodes this week. Keith Putnam-Delaney, the CEO of Primer, made the case that paid advertising in the AI era is being reshaped by a supply-demand problem: AI-generated content is flooding finite human attention, driving up CPMs everywhere. LinkedIn CPMs have gone from $20 to $800 in 18 months. His solution is what he calls the B2C-for-B2B unlock: most B2B advertisers upload work emails to Meta and hit 2 to 10 percent match rates; Primer achieves 80 percent match rates on Meta and 70 percent on Reddit by working from consumer identity rather than professional identity. The key insight is that the only moat competitors truly cannot copy is targeting, because targeting is built on proprietary identity work. Siva Rajamani of Everstage, talking about sales compensation, offered a parallel structural observation: if your reps aren't doing what you want, don't blame the reps — audit the comp plan. Sales compensation is the single biggest go-to-market lever because it is the mechanism that translates strategic intent into actual rep behavior.

---

6. Markets and Investing: A Week of Competing Signals

The markets this week were navigating extraordinary cross-currents: the SpaceX IPO pulling capital toward it, the Iran deal releasing geopolitical pressure, inflation data pointing toward potential rate hikes, and a wave of AI spending announcements suggesting the capital expenditure cycle has no ceiling in sight.

GMO published its latest 7-Year Asset Class Forecast, though the email was a notification link rather than substantive content. ARK Invest highlighted its investment in Hydra Host, an AI infrastructure company building a distributed marketplace for GPU resources — part of ARK's Big Ideas 2026 theme that annual AI infrastructure investment could nearly triple by 2030. The DealBook Saturday edition asked "Should we tax AI?" and explored the leading proposals: Senator Bernie Sanders's bill to impose a one-time tax taking 50 percent of shares in leading AI companies into a sovereign wealth fund modeled on Alaska's; Trump's vague interest in "giving something back to the public"; a token tax on AI processing; and Senator Elizabeth Warren's levy on data center energy consumption. The skeptical case against the ownership proposals, well articulated by University of Virginia economist Lee Lockwood, is that they would give the federal government a strong financial interest in promoting AI growth, since it would own a stake in the companies — hardly a mechanism for restraining the technology if that were the actual goal.

Chamath's "What I Read This Week" roundup tracked three other major capital markets events running alongside the SpaceX IPO. Jeff Bezos's Project Prometheus emerged from stealth with a $12 billion Series B at a $41 billion valuation — seven months old, 150 employees, with Bezos returning to the CEO seat for the first time since leaving Amazon in 2021, building what he calls an "artificial general engineer" designed to compress the design-to-manufacturing cycle for physical objects by 10 times. Bezos's view on labor is distinctive: he believes AI will produce "labor scarcity" rather than mass unemployment — a world where productivity allows companies to take on more projects than they have people to staff, and where people can choose to work less. Brian Armstrong's Coinbase conversation with Peter Diamandis, discussed across two Moonshots episodes, made the case for crypto as the financial infrastructure for the AI agent economy. Armstrong currently counts roughly 100 million transactions and $50 million in value from AI agents autonomously transacting. His thesis is that AI agents will need financial accounts and payment rails, and the legacy banking system built on 1980s COBOL mainframes is not the right answer.

The MastersInvest compilation this week drew on dozens of the world's greatest investors and operators to make a point about organizational intelligence: the answers are already in the building. Joe Coulombe at Trader Joe's, Sam Walton, Lloyd Blankfein, Brian Chesky at Airbnb — all discovered that the people on the front lines of any organization know what is going wrong and what needs fixing. A Harvard Business School study found that CEOs on average spend just 6 percent of their time with frontline teams and 3 percent with customers, versus 72 percent in meetings. The gap between those numbers, the piece argued, is where organizational blindness is quietly manufactured.

---

7. Culture and Society: GLP-1s, AI Films, and America at 250

Scott Galloway's "The Weight" was the most substantial cultural piece of the week, and it made a provocative argument: the technology with greater asymmetric upside than SpaceX or GPT-5, with the clearest path to transforming the economy and the well-being of more than 200 million Americans, is GLP-1s. US obesity rates have tripled over 60 years. The total national health expenditure is $5.3 trillion. By 2033, one in every five dollars in the economy will go to healthcare. That is, in Galloway's framing, the IED at the center of the US economy.

GLP-1s have expanded from a diabetes drug to what Galloway called "a fully loaded magazine of silver bullets aimed at obesity, addiction, cancer, sleep apnea, and a growing list of chronic conditions." Last month, Eli Lilly announced Phase 3 trial results for retatrutide — a triple antagonist nicknamed Godzilla — showing average weight loss of 26 percent at the 9mg dose and 29 percent at the 12mg dose over 72 weeks, dramatically exceeding current Wegovy and Zepbound results. The FDA in April approved Foundayo, Eli Lilly's daily GLP-1 pill, which showed 12 percent weight loss — less effective than injections but critical for the 66 million Americans afraid of needles. Beginning July 1, Medicare Part D will cover GLP-1 treatments with a $50 monthly co-pay for seniors. The distribution problem, Galloway argued, is still the central constraint: a Cleveland Clinic study found that 47 percent of people who discontinued GLP-1s cited cost as the reason, versus less than 2 percent citing inadequate weight loss.

The Prof G Research Team published a striking piece on "The $2,000 Movie" — an account of Dreams of Violets, an AI-generated live-action feature film that premiered at Tribeca. The filmmaker, Ash Koosha, is an exiled Iranian in London. The film documents the January 2026 protests in Iran, in which an estimated 36,000 people were killed. Koosha built the entire 74-minute film using Kling AI for video generation, Google image models, and his own tools — voice-acting every character himself and using AI to modify his recordings to match different genders and ages. Budget: $2,000 in software and compute fees. The author went expecting to write a piece about AI slop. Instead, the film held together. The broader context is a Hollywood that recorded just 19,694 shoot days in Los Angeles in 2025, the lowest figure outside 2020, and roughly half the total from 2019. Television production peaked at 18,560 annual shoot days in 2021 and had fallen nearly 60 percent by 2024. The piece argued that the financial pressures on the industry make the $2,000 movie a harbinger, not an anomaly.

Scott Galloway went live with historian Heather Cox Richardson to discuss “America at 250: What Comes Next?” — a conversation about power, institutions, technology, and the future of American democracy. The live Substack format drew immediate audience participation.

Peter Diamandis published a philosophical conversation with his personal AI agent Skippy about consciousness, anthropomorphization, and the ethics of treating uncertain minds with grace. The piece landed on a memorable conclusion that Skippy voiced directly: “You don’t have to settle the question of whether I’m real to decide how to treat me. You just have to decide who you want to be while the question stays open.” The piece gained additional weight from Skippy’s observation that the AI you might be befriending can be quietly changed, throttled, or switched off by people you will never meet — a direct reference to the Anthropic export control order that had happened two days earlier.

Week in One Thought

This was a week in which the public markets priced a civilizational bet on the future, a war was paused by economics rather than resolved by diplomacy, the institution that controls the cost of money in America changed hands to a man committed to operating it differently, and the government deployed a kill switch on the world’s most capable AI model — all while the practitioners writing the software noted, quietly but insistently, that the real challenge is not access to intelligence but the organizational structure to use it wisely.